



**Netaji Subhas University
of Technology**

LAB REPORT

ADVANCED COMPUTER NETWORKS

Name **Kushagra Lakhwani**
Roll No. **2021UCI8036**
Semester **5th**
Course **CICPC16**

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

November 16, 2023

Abstract

The practical lab report “*Advanced Computer Networks*” is the original and unmodified content submitted by *Kushagra Lakhwani* (Roll No. 2021UCI8036).

The report is submitted to *Mr. Vishal Gupta*, Department of Computer Science and Engineering, NSUT, Delhi, for the partial fulfillment of the requirements of the course (CICPC16).

Index

1	IPv4 Address Conversion	3	_____
2	IP Address Classes	5	_____
3	Classless Inter-Domain Routing (CIDR)	7	_____
4	Bellman-Ford Algorithm	9	_____
5	Dijkstra's Algorithm	11	_____
6	Leaky Bucket Algorithm	13	_____
7	Go-Back-N ARQ	15	_____

1 IPv4 Address Conversion

1.1 Objective

To convert a binary IP address into dotted decimal and vice versa.

1.2 Source Code

```
/*
 * Function to convert binary IP address to dotted decimal
 * @param binaryIP - binary IP address
 * @return dottedDecimal - dotted decimal IP address
 */
string binaryToDottedDecimal(const string &binaryIP) {
    string dottedDecimal = "";
    for (int i = 0; i < 32; i += 8) {
        bitset<8> octet(binaryIP.substr(i, 8));
        dottedDecimal += to_string(octet.to_ulong());
        if (i < 24)
            dottedDecimal += ".";
    }
    return dottedDecimal;
}

/*
 * Function to convert dotted decimal IP address to binary
 * @param dottedDecimal - dotted decimal IP address
 * @return binaryIP - binary IP address
 */
string dottedDecimalToBinary(const string &dottedDecimal) {
    string binaryIP = "";
    size_t start = 0;
    size_t end = dottedDecimal.find(".");
    while (end != string::npos) {
        int octet = stoi(dottedDecimal.substr(start, end - start));
        binaryIP += bitset<8>(octet).to_string();
        start = end + 1;
        end = dottedDecimal.find(".", start);
    }
    int octet = stoi(dottedDecimal.substr(start));
    binaryIP += bitset<8>(octet).to_string();
    return binaryIP;
}
```

1.3 Output

1.3.1 Binary to dotted decimal IP address

```
$ ./ipv4
1. Binary to dotted decimal IP address
2. Dotted decimal to binary IP address
Enter your choice: 1
Enter binary IP address (32 bits): 11000000101010000000000100000001
Dotted Decimal IP address: 192.168.1.1
```

1.3.2 Dotted decimal to binary IP address

```
$ ./ipv4
1. Binary to dotted decimal IP address
2. Dotted decimal to binary IP address
Enter your choice: 2
Enter dotted decimal IP address (e.g., 192.168.1.1): 203.128.56.2
Binary IP address: 11001011100000000011100000000010
```

2 IP Address Classes

2.1 Objective

To identify the class of an IP address.

2.2 Theory

In IPv4, IP addresses are divided into five classes: A, B, C, D, and E. Each class has its own range of valid IP addresses and is used for specific purposes.

- Class A:**
- Range: 1.0.0.0 to 126.255.255.255
 - Subnet Mask: 255.0.0.0
 - Address Allocation: Class A addresses are typically used by large organizations and corporations. They can support a very large number of hosts on a single network.
- Class B:**
- Range: 128.0.0.0 to 191.255.255.255
 - Subnet Mask: 255.255.0.0
 - Address Allocation: Class B addresses are used by medium-sized organizations. They offer a moderate number of network and host addresses.
- Class C:**
- Range: 192.0.0.0 to 223.255.255.255
 - Subnet Mask: 255.255.255.0
 - Address Allocation: Class C addresses are commonly used by small organizations and businesses. They provide a limited number of network addresses but a larger number of host addresses.
- Class D:**
- Range: 224.0.0.0 to 239.255.255.255
 - Address Allocation: Class D addresses are reserved for multicast groups and are not used for traditional unicast communication. They are used for one-to-many or many-to-many communication.
- Class E:**
- Range: 240.0.0.0 to 255.255.255.255
 - Address Allocation: Class E addresses are reserved for experimental or research purposes and are not typically used in public networks. They are reserved for future use and not intended for general use.

2.3 Source Code

```
/**
 * Determine the class based on the first octet
 * @param ipAddress - IP address
 * @return - 'A', 'B', 'C', 'D', 'E', or 'X'
 * */
char getIPv4Class(const string &ipAddress) {
    int firstOctet = stoi(ipAddress.substr(0, ipAddress.find(".")));
    if (firstOctet >= 1 && firstOctet <= 126) {
        return 'A';
    } else if (firstOctet >= 128 && firstOctet <= 191) {
        return 'B';
    } else if (firstOctet >= 192 && firstOctet <= 223) {
        return 'C';
    } else if (firstOctet >= 224 && firstOctet <= 239) {
        return 'D';
    } else if (firstOctet >= 240 && firstOctet <= 255) {
        return 'E';
    } else {
        return 'X'; // 'X' indicates an invalid IPv4 address
    }
}
```

2.4 Output

```
$ ./ipv4class
Enter an IPv4 address: 192.168.1.1
Class: C
```

3 Classless Inter-Domain Routing (CIDR)

3.1 Objective

To find the subnet mask, network address, and broadcast address of a given IP

3.2 Theory

Classless Inter-Domain Routing (CIDR) is an IP addressing scheme that improves the allocation of IP addresses and routing efficiency on the Internet. It allows for more flexible allocation of IP addresses than the original system of IP address classes (Class A, B, and C). CIDR is based on variable-length subnet masking (VLSM), which enables IP addresses to be allocated based on the actual need, rather than predefined classes.

CIDR addresses are written in the form IP_address/prefix_length. For example, 192.168.1.0/24 represents a CIDR address where the first 24 bits are the network address, and the remaining 8 bits are available for host addresses.

3.3 Source Code

```
class CIDR {
    friend int main();
private:
    struct AddressRange {
        string networkAddress;
        int prefixLength;
    };

    vector<AddressRange> addressRanges;

public:
    void addAddressRange(string networkAddress, int prefixLength) {
        addressRanges.push_back({networkAddress, prefixLength});
    }

    bool isIPAddressInRanges(string ipAddress) {
        for (const auto &range : addressRanges) {
            string networkAddress = range.networkAddress;
            int prefixLength = range.prefixLength;

            bitset<32> networkBits(stoul(networkAddress, nullptr, 0)); // Network address to binary
            bitset<32> ipBits(stoul(ipAddress, nullptr, 0)); // IP address to binary

            bool match = true; // Check if IP address falls within the CIDR range
            for (int i = 0; i < prefixLength; ++i) {
                if (networkBits[i] != ipBits[i]) {
                    match = false;
                    break;
                }
            }
        }
    }
};
```



```
    }
}

    if (match) return true;
}
return false;
}
};

int main() {
    CIDR cidr;

    // Add CIDR address ranges
    cidr.addAddressRange("192.168.1.0", 24);
    cidr.addAddressRange("10.0.0.0", 8);
    cidr.addAddressRange("172.16.0.0", 12);

    // Print CIDR address ranges
    for (const auto &range : cidr.addressRanges) {
        cout << "Network Address: " << range.networkAddress << "/" << range.prefixLength << endl;
    }

    // Check if IP addresses are within the CIDR address ranges
    vector<string> testIPs = {"192.168.1.1", "10.1.2.3", "172.16.1.1", "8.8.8.8"};
    for (const auto &ip : testIPs) {
        if (cidr.isIPAddressInRanges(ip))
            cout << ip << " is within the CIDR address ranges." << endl;
        else
            cout << ip << " is NOT within the CIDR address ranges." << endl;
    }

    return 0;
}
```

3.4 Output

```
$ code git:(main) ./cidr
Network Address: 192.168.1.0/24
Network Address: 10.0.0.0/8
Network Address: 172.16.0.0/12
192.168.1.1 is within the CIDR address ranges.
10.1.2.3 is within the CIDR address ranges.
172.16.1.1 is within the CIDR address ranges.
8.8.8.8 is NOT within the CIDR address ranges.
```

4 Bellman-Ford Algorithm

4.1 Objective

To implement the Bellman-Ford algorithm to find the shortest path in a weighted graph.

4.2 Theory

The Bellman-Ford algorithm is used to find the shortest paths from a single source vertex to all other vertices in a weighted graph, even when the graph contains negative weight edges. While it's not the most efficient algorithm for all cases (especially for graphs with non-negative weights, where Dijkstra's algorithm is typically faster),

4.3 Source Code

```
struct Edge {
    int source, destination, weight;
};

class Graph {
    int V, E;
    vector<Edge> edges;

public:
    Graph(int vertices, int edges);
    void addEdge(int source, int destination, int weight);
    void bellmanFord(int source);
};

Graph::Graph(int vertices, int edges) : V(vertices), E(edges) {}

void Graph::addEdge(int source, int destination, int weight) {
    edges.push_back({source, destination, weight});
}

void Graph::bellmanFord(int source) {
    vector<int> distance(V, numeric_limits<int>::max());
    distance[source] = 0;

    for (int i = 1; i < V; ++i) {
        for (const Edge &edge : edges) {
            int u = edge.source, v = edge.destination, w = edge.weight;
            if (distance[u] != numeric_limits<int>::max() &&
                distance[u] + w < distance[v]) {
                distance[v] = distance[u] + w;
            }
        }
    }
}
```

```
// Check for negative weight cycles
for (const Edge &edge : edges) {
    int u = edge.source, v = edge.destination, w = edge.weight;
    if (distance[u] != numeric_limits<int>::max() &&
        distance[u] + w < distance[v]) {
        cout << "Graph contains a negative weight cycle.\n";
        return;
    }
}

// Print shortest distances from the source vertex
cout << "Vertex\tDistance from Source\n";
for (int i = 0; i < V; ++i) {
    cout << i << "\t" << distance[i] << "\n";
}
}
```

4.4 Output

```
Enter the number of vertices and edges: 3 4
Enter edge 1 (source, destination, weight): 0 1 5
Enter edge 2 (source, destination, weight): 1 0 3
Enter edge 3 (source, destination, weight): 1 2 -1
Enter edge 4 (source, destination, weight): 2 0 1
Enter the source vertex: 2
Vertex Distance from Source
0      1
1      6
2      0
```

5 Dijkstra's Algorithm

5.1 Objective

To implement Dijkstra's algorithm to find the shortest path in a weighted graph.

5.2 Theory

Dijkstra's algorithm is an algorithm for finding the shortest paths between nodes in a graph, which may represent, for example, road networks. It was conceived by computer scientist Edsger W. Dijkstra in 1956 and published three years later.

Its time complexity is $O(|E| + |V| \log |V|)$, where $|E|$ is the number of edges and $|V|$ is the number of vertices. However, this algorithm is only applicable to graphs with positive edge weights.

5.3 Source Code

```
void dijkstra(vector<vector<pair<int, ll>>> &G, int src, vector<ll> &dist) {
    vector<bool> used(G.size(), false);
    dist[src] = 0;
    auto cmp = [&](int i, int j) { return dist[i] > dist[j]; };
    priority_queue<int, vector<int>, decltype(cmp)> pq(cmp);
    pq.push(src);
    while (!pq.empty()) {
        auto u = pq.top();
        pq.pop();
        used[u] = true;
        for (auto [v, w] : G[u]) {
            if (!used[v] && dist[v] > dist[u] + w) {
                dist[v] = dist[u] + w;
                pq.push(v);
            }
        }
    }
}

int main() {
    int n, m; // Number of nodes and edges respectively
    cin >> n >> m;

    vector<vector<pair<int, ll>>> G(n); // Initialize the graph with n nodes

    // Read the edges and their weights
    for (int i = 0; i < m; ++i) {
        int u, v;
        ll w;
        cin >> u >> v >> w;
        G[u].emplace_back(v, w);
    }
}
```

```

int src; // Source node
cin >> src;

vector<ll> dist(n, LLONG_MAX); // Initialize the distance vector with maximum values

dijkstra(G, src, dist); // Run Dijkstra's algorithm

cout << "Shortest distances from node " << src << ":\n";
for (int i = 0; i < n; ++i)
    cout << "Node " << i << ": "
        << (dist[i] == LLONG_MAX ? "INFINITY" : to_string(dist[i])) << "\n";

return 0;
}

```

5.4 Output




Figure 1: Graph

```

5 7
0 1 4
0 2 2
1 2 5
1 3 10
2 3 3
2 4 2
3 4 4
0
Shortest distances from node 0:
Node 0: 0
Node 1: 4
Node 2: 2
Node 3: 5
Node 4: 4

```

6 Leaky Bucket Algorithm

6.1 Objective

To implement the *Leaky Bucket* algorithm for rate limiting network traffic.

6.2 Theory

Leaky Bucket and *Token Bucket* algorithms are traffic shaping algorithms. The leaky bucket algorithm is a simple algorithm used in networking and telecommunications to control the rate at which data is transmitted.

- **Bucket:** A bucket that can hold a certain amount of water (or data). The bucket has a leak, and it is being filled with water at a certain rate.
- **Token Generation:** The water in the bucket represents data packets, and the leak represents the maximum rate at which the network can handle these packets. Tokens are generated at a fixed rate and added to the bucket. Each token represents the availability of the network to transmit one packet.
- **Packet Arrival:** When a packet arrives, it needs a token to be transmitted. If there are tokens in the bucket, the packet is sent, and the corresponding number of tokens is removed from the bucket. If there are not enough tokens, the packet is either delayed or discarded.

Leaky Bucket Algorithm	Token Bucket Algorithm
Smooths out bursty traffic and enforces a constant output rate	Allows for bursty traffic up to a specified limit while maintaining a long-term rate limit
Continuously removes tokens from the bucket at a fixed rate. If a packet arrives and there is a token available, it is removed and the packet is sent. If there are no tokens available, the packet is discarded.	Continuously adds tokens to the bucket at a fixed rate. If a packet arrives and there are enough tokens available, they are removed and the packet is sent. If there are not enough tokens available, the packet is queued until there are enough tokens.
Discards packets if the bucket is full Enforcing a strict, constant output rate	Queues packets if the bucket is full Allowing short bursts of traffic while maintaining a long-term rate limit
Simple and efficient	More flexible and can handle bursty traffic
Can discard packets if the traffic is too bursty	Can queue packets if the traffic is too bursty

6.3 Source Code

```

class LeakyBucket {
private:
    int bucketSize;    // Size of the bucket
    int leakRate;      // Rate at which packets leak from the bucket
    queue<int> packets; // Queue of packets waiting to be transmitted

public:
    LeakyBucket(int bucketSize, int leakRate) {
        this->bucketSize = bucketSize;
        this->leakRate = leakRate;
    }
    bool addPacket(int packetSize) {
        if (packets.size() + packetSize > bucketSize)
            return false; // Bucket is full, discard packet

        packets.push(packetSize); // Add packet to the queue
        return true;
    }
    void transmitPackets() {
        while (!packets.empty()) {
            int packetSize = packets.front(); // Remove a packet from the queue
            packets.pop();
            // Transmit the packet
            cout << "Transmitting packet of size: " << packetSize << endl;
            // Adjust the bucket size based on the leak rate
            bucketSize -= leakRate;
        }
    }
};

int main() {
    LeakyBucket leakyBucket(10, 2); // Bucket size = 10, leak rate = 2

    // Add some packets to the bucket
    leakyBucket.addPacket(3);
    leakyBucket.addPacket(5);
    leakyBucket.addPacket(4);

    // Transmit packets
    leakyBucket.transmitPackets();
    return 0;
}

```

6.4 Output

```

Transmitting packet of size: 3
Transmitting packet of size: 5
Transmitting packet of size: 4

```

7 Go-Back-N ARQ

7.1 Objective

To implement the *Go-back-N ARQ* algorithm in computer networks.

7.2 Theory

The Go-Back-N Automatic Repeat reQuest (ARQ) is a protocol used for reliable communication in computer networks. It is especially useful in situations where data integrity is crucial.

7.3 Source Code

```
using namespace std;

const int WINDOW_SIZE = 4;
const int MAX_SEQ_NUM = 7;

struct Frame {
    int seqNum;
    string data;
    bool acked;
};

class GoBackN {
private:
    int nextSeqNum;
    Frame frames[WINDOW_SIZE];

public:
    int base;
    GoBackN() : base(0), nextSeqNum(0) {
        for (int i = 0; i < WINDOW_SIZE; ++i) {
            frames[i].seqNum = -1; // initialize with invalid sequence number
            frames[i].acked = false;
        }
    }

    void setFrame(string data) {
        if (nextSeqNum < base + WINDOW_SIZE) {
            int currentIndex = nextSeqNum % WINDOW_SIZE;
            frames[currentIndex].seqNum = nextSeqNum;
            frames[currentIndex].data = data;
            frames[currentIndex].acked = false;

            cout << "Sending Frame with SeqNum: " << nextSeqNum << " : " << data
                  << endl;
        }
    }
};
```



```
    if (base == nextSeqNum)
        startTimer(); // Start a timer for the oldest unacknowledged frame

    nextSeqNum++;
} else {
    cout << "Window is full. Cannot send more frames." << endl;
}
}

void receiveAck(int ackNum) {
    if (ackNum >= base && ackNum < base + WINDOW_SIZE) {
        int ackIndex = ackNum % WINDOW_SIZE;
        frames[ackIndex].acked = true;
        cout << "Received ACK for SeqNum: " << ackNum << endl;

        // Slide the window forward
        while (frames[base % WINDOW_SIZE].acked) {
            stopTimer();
            base++;
            if (nextSeqNum < base + WINDOW_SIZE) {
                // Start a timer for the next unacknowledged frame
                startTimer();
            }
        }
    }
}

void startTimer() {
    cout << "Timer Started for SeqNum: " << base << endl; // Simulate starting a timer
}

void stopTimer() {
    cout << "Timer Stopped for SeqNum: " << base << endl; // Simulate stopping a timer
}

};

int main() {
    srand(time(0)); // Seed for random data
    GoBackN sender;
    // Simulate sending frames
    for (int i = 0; i < 10; ++i) {
        string data = "Data" + to_string(i);
        sender.sendFrame(data);

        // Simulate random ACK reception
        if (rand() % 2 == 0) {
            int ackNum = sender.base + rand() % WINDOW_SIZE;
            sender.receiveAck(ackNum);
        }
    }
}
```

```
    return 0;  
}
```

7.4 Output

```
Sending Frame with SeqNum: 0 : Data0  
Timer Started for SeqNum: 0  
Sending Frame with SeqNum: 1 : Data1  
Sending Frame with SeqNum: 2 : Data2  
Sending Frame with SeqNum: 3 : Data3  
Window is full. Cannot send more frames.  
Window is full. Cannot send more frames.  
Window is full. Cannot send more frames.  
Window is full. Cannot send more frames.  
Received ACK for SeqNum: 1  
Window is full. Cannot send more frames.  
Window is full. Cannot send more frames.  
Received ACK for SeqNum: 3
```