



**Netaji Subhas University
of Technology**

LAB REPORT

ADVANCED COMPUTER NETWORKS

Name **Kushagra Lakhwani**
Roll No. **2021UCI8036**
Semester **5th**
Course **CICPC16**

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

October 11, 2023

Abstract

The practical lab report “*Advanced Computer Networks*” is the original and unmodified content submitted by *Kushagra Lakhwani* (Roll No. 2021UCI8036).

The report is submitted to *Mr. Vishal Gupta*, Department of Computer Science and Engineering, NSUT, Delhi, for the partial fulfillment of the requirements of the course (CICPC16).

Index

1	IPv4 Address Conversion	3	_____
2	IP Address Classes	5	_____
3	Classless Inter-Domain Routing (CIDR)	7	_____
4	Distance Vector Routing Algorithm	9	_____
5	Bellman-Ford Algorithm	11	_____
6	Dijkstra's Algorithm	13	_____

1 IPv4 Address Conversion

1.1 Objective

To convert a binary IP address into dotted decimal and vice versa.

1.2 Source Code

```
/*
 * Function to convert binary IP address to dotted decimal
 * @param binaryIP - binary IP address
 * @return dottedDecimal - dotted decimal IP address
 */
string binaryToDottedDecimal(const string &binaryIP) {
    string dottedDecimal = "";
    for (int i = 0; i < 32; i += 8) {
        bitset<8> octet(binaryIP.substr(i, 8));
        dottedDecimal += to_string(octet.to_ulong());
        if (i < 24)
            dottedDecimal += ".";
    }
    return dottedDecimal;
}

/*
 * Function to convert dotted decimal IP address to binary
 * @param dottedDecimal - dotted decimal IP address
 * @return binaryIP - binary IP address
 */
string dottedDecimalToBinary(const string &dottedDecimal) {
    string binaryIP = "";
    size_t start = 0;
    size_t end = dottedDecimal.find(".");
    while (end != string::npos) {
        int octet = stoi(dottedDecimal.substr(start, end - start));
        binaryIP += bitset<8>(octet).to_string();
        start = end + 1;
        end = dottedDecimal.find(".", start);
    }
    int octet = stoi(dottedDecimal.substr(start));
    binaryIP += bitset<8>(octet).to_string();
    return binaryIP;
}
```

1.3 Output

1.3.1 Binary to dotted decimal IP address

```
$ ./ipv4
1. Binary to dotted decimal IP address
2. Dotted decimal to binary IP address
Enter your choice: 1
Enter binary IP address (32 bits): 11000000101010000000000100000001
Dotted Decimal IP address: 192.168.1.1
```

1.3.2 Dotted decimal to binary IP address

```
$ ./ipv4
1. Binary to dotted decimal IP address
2. Dotted decimal to binary IP address
Enter your choice: 2
Enter dotted decimal IP address (e.g., 192.168.1.1): 203.128.56.2
Binary IP address: 11001011100000000011100000000010
```

2 IP Address Classes

2.1 Objective

To identify the class of an IP address.

2.2 Theory

In IPv4, IP addresses are divided into five classes: A, B, C, D, and E. Each class has its own range of valid IP addresses and is used for specific purposes.

- Class A:**
- Range: 1.0.0.0 to 126.255.255.255
 - Subnet Mask: 255.0.0.0
 - Address Allocation: Class A addresses are typically used by large organizations and corporations. They can support a very large number of hosts on a single network.
- Class B:**
- Range: 128.0.0.0 to 191.255.255.255
 - Subnet Mask: 255.255.0.0
 - Address Allocation: Class B addresses are used by medium-sized organizations. They offer a moderate number of network and host addresses.
- Class C:**
- Range: 192.0.0.0 to 223.255.255.255
 - Subnet Mask: 255.255.255.0
 - Address Allocation: Class C addresses are commonly used by small organizations and businesses. They provide a limited number of network addresses but a larger number of host addresses.
- Class D:**
- Range: 224.0.0.0 to 239.255.255.255
 - Address Allocation: Class D addresses are reserved for multicast groups and are not used for traditional unicast communication. They are used for one-to-many or many-to-many communication.
- Class E:**
- Range: 240.0.0.0 to 255.255.255.255
 - Address Allocation: Class E addresses are reserved for experimental or research purposes and are not typically used in public networks. They are reserved for future use and not intended for general use.

2.3 Source Code

```
/**
 * Determine the class based on the first octet
 * @param ipAddress - IP address
 * @return - 'A', 'B', 'C', 'D', 'E', or 'X'
 * */
char getIPv4Class(const string &ipAddress) {
    int firstOctet = stoi(ipAddress.substr(0, ipAddress.find(".")));
    if (firstOctet >= 1 && firstOctet <= 126) {
        return 'A';
    } else if (firstOctet >= 128 && firstOctet <= 191) {
        return 'B';
    } else if (firstOctet >= 192 && firstOctet <= 223) {
        return 'C';
    } else if (firstOctet >= 224 && firstOctet <= 239) {
        return 'D';
    } else if (firstOctet >= 240 && firstOctet <= 255) {
        return 'E';
    } else {
        return 'X'; // 'X' indicates an invalid IPv4 address
    }
}
```

2.4 Output

```
$ ./ipv4class
Enter an IPv4 address: 192.168.1.1
Class: C
```

3 Classless Inter-Domain Routing (CIDR)

3.1 Objective

To find the subnet mask, network address, and broadcast address of a given IP

3.2 Theory

Classless Inter-Domain Routing (CIDR) is an IP addressing scheme that improves the allocation of IP addresses and routing efficiency on the Internet. It allows for more flexible allocation of IP addresses than the original system of IP address classes (Class A, B, and C). CIDR is based on variable-length subnet masking (VLSM), which enables IP addresses to be allocated based on the actual need, rather than predefined classes.

CIDR addresses are written in the form IP_address/prefix_length. For example, 192.168.1.0/24 represents a CIDR address where the first 24 bits are the network address, and the remaining 8 bits are available for host addresses.

3.3 Source Code

```
class CIDR {
    friend int main();
private:
    struct AddressRange {
        string networkAddress;
        int prefixLength;
    };

    vector<AddressRange> addressRanges;

public:
    void addAddressRange(string networkAddress, int prefixLength) {
        addressRanges.push_back({networkAddress, prefixLength});
    }

    bool isIPAddressInRanges(string ipAddress) {
        for (const auto &range : addressRanges) {
            string networkAddress = range.networkAddress;
            int prefixLength = range.prefixLength;

            bitset<32> networkBits(stoul(networkAddress, nullptr, 0)); // Network address to binary
            bitset<32> ipBits(stoul(ipAddress, nullptr, 0)); // IP address to binary

            bool match = true; // Check if IP address falls within the CIDR range
            for (int i = 0; i < prefixLength; ++i) {
                if (networkBits[i] != ipBits[i]) {
                    match = false;
                    break;
                }
            }
        }
    }
};
```



```
    }  
}  
  
    if (match) return true;  
}  
    return false;  
}  
};  
  
int main() {  
    CIDR cidr;  
  
    // Add CIDR address ranges  
    cidr.addAddressRange("192.168.1.0", 24);  
    cidr.addAddressRange("10.0.0.0", 8);  
    cidr.addAddressRange("172.16.0.0", 12);  
  
    // Print CIDR address ranges  
    for (const auto &range : cidr.addressRanges) {  
        cout << "Network Address: " << range.networkAddress << "/" << range.prefixLength << endl;  
    }  
  
    // Check if IP addresses are within the CIDR address ranges  
    vector<string> testIPs = {"192.168.1.1", "10.1.2.3", "172.16.1.1", "8.8.8.8"};  
    for (const auto &ip : testIPs) {  
        if (cidr.isIPAddressInRanges(ip))  
            cout << ip << " is within the CIDR address ranges." << endl;  
        else  
            cout << ip << " is NOT within the CIDR address ranges." << endl;  
    }  
  
    return 0;  
}
```

3.4 Output

```
$ code git:(main) ./cidr  
Network Address: 192.168.1.0/24  
Network Address: 10.0.0.0/8  
Network Address: 172.16.0.0/12  
192.168.1.1 is within the CIDR address ranges.  
10.1.2.3 is within the CIDR address ranges.  
172.16.1.1 is within the CIDR address ranges.  
8.8.8.8 is NOT within the CIDR address ranges.
```

4 Distance Vector Routing Algorithm

4.1 Objective

To implement the Distance-Vector Routing algorithm.

4.2 Theory

A distance-vector routing (DVR) protocol requires that a router inform its neighbors of topology changes periodically. Historically known as the old *ARPANET* routing algorithm or known as *Bellman-Ford* algorithm.

Each router maintains a Distance Vector table containing the distance between itself and *all* possible destination nodes. Distances, based on a chosen metric, are computed using information from the neighbors' distance vectors.

4.3 Information Kept by DV Router

- Each router has an ID.
- Associated with each link connected to a router, there is a link cost (static or dynamic).
- Intermediate hops.

4.4 Distance Vector Table Initialization

- Distance to itself = 0
- Distance to all other routers = infinity number.

4.5 Distance Vector Algorithm

1. A router transmits its distance vector to each of its neighbors in a routing packet.
2. Each router receives and saves the most recently received distance vector from each of its neighbors.
3. A router recalculates its distance vector when:
 - It receives a distance vector from a neighbor containing different information than before.
 - It discovers that a link to a neighbor has gone down.

4. The DV calculation is based on minimizing the cost to each destination:

$Dx(y)$ = Estimate of least cost from x to y

$C(x, v)$ = Node x knows cost to each neighbor v

$Dx = [Dx(y) : y \in N]$ = Node x maintains distance vector

Node x also maintains its neighbors' distance vectors:

For each neighbor v , x maintains $Dv = [Dv(y) : y \in N]$

5 Bellman-Ford Algorithm

5.1 Objective

To implement the Bellman-Ford algorithm to find the shortest path in a weighted graph.

5.2 Theory

The Bellman-Ford algorithm is used to find the shortest paths from a single source vertex to all other vertices in a weighted graph, even when the graph contains negative weight edges. While it's not the most efficient algorithm for all cases (especially for graphs with non-negative weights, where Dijkstra's algorithm is typically faster),

5.3 Source Code

```
struct Edge {
    int source, destination, weight;
};

class Graph {
    int V, E;
    vector<Edge> edges;

public:
    Graph(int vertices, int edges);
    void addEdge(int source, int destination, int weight);
    void bellmanFord(int source);
};

Graph::Graph(int vertices, int edges) : V(vertices), E(edges) {}

void Graph::addEdge(int source, int destination, int weight) {
    edges.push_back({source, destination, weight});
}

void Graph::bellmanFord(int source) {
    vector<int> distance(V, numeric_limits<int>::max());
    distance[source] = 0;

    for (int i = 1; i < V; ++i) {
        for (const Edge &edge : edges) {
            int u = edge.source, v = edge.destination, w = edge.weight;
            if (distance[u] != numeric_limits<int>::max() &&
                distance[u] + w < distance[v]) {
                distance[v] = distance[u] + w;
            }
        }
    }
}
```

```
// Check for negative weight cycles
for (const Edge &edge : edges) {
    int u = edge.source, v = edge.destination, w = edge.weight;
    if (distance[u] != numeric_limits<int>::max() &&
        distance[u] + w < distance[v]) {
        cout << "Graph contains a negative weight cycle.\n";
        return;
    }
}

// Print shortest distances from the source vertex
cout << "Vertex\tDistance from Source\n";
for (int i = 0; i < V; ++i) {
    cout << i << "\t" << distance[i] << "\n";
}
}
```

5.4 Output

```
Enter the number of vertices and edges: 3 4
Enter edge 1 (source, destination, weight): 0 1 5
Enter edge 2 (source, destination, weight): 1 0 3
Enter edge 3 (source, destination, weight): 1 2 -1
Enter edge 4 (source, destination, weight): 2 0 1
Enter the source vertex: 2
Vertex Distance from Source
0      1
1      6
2      0
```

6 Dijkstra's Algorithm

6.1 Objective

To implement Dijkstra's algorithm to find the shortest path in a weighted graph.

6.2 Theory

Dijkstra's algorithm is an algorithm for finding the shortest paths between nodes in a graph, which may represent, for example, road networks. It was conceived by computer scientist Edsger W. Dijkstra in 1956 and published three years later.

Its time complexity is $O(|E| + |V| \log |V|)$, where $|E|$ is the number of edges and $|V|$ is the number of vertices. However, this algorithm is only applicable to graphs with positive edge weights.

6.3 Source Code

```
void dijkstra(vector<vector<pair<int, ll>>> &G, int src, vector<ll> &dist) {
    vector<bool> used(G.size(), false);
    dist[src] = 0;
    auto cmp = [&](int i, int j) { return dist[i] > dist[j]; };
    priority_queue<int, vector<int>, decltype(cmp)> pq(cmp);
    pq.push(src);
    while (!pq.empty()) {
        auto u = pq.top();
        pq.pop();
        used[u] = true;
        for (auto [v, w] : G[u]) {
            if (!used[v] && dist[v] > dist[u] + w) {
                dist[v] = dist[u] + w;
                pq.push(v);
            }
        }
    }
}

int main() {
    int n, m; // Number of nodes and edges respectively
    cin >> n >> m;

    vector<vector<pair<int, ll>>> G(n); // Initialize the graph with n nodes

    // Read the edges and their weights
    for (int i = 0; i < m; ++i) {
        int u, v;
        ll w;
        cin >> u >> v >> w;
        G[u].emplace_back(v, w);
    }
}
```

```

int src; // Source node
cin >> src;

vector<ll> dist(n, LLONG_MAX); // Initialize the distance vector with maximum values

dijkstra(G, src, dist); // Run Dijkstra's algorithm

cout << "Shortest distances from node " << src << ":\n";
for (int i = 0; i < n; ++i)
    cout << "Node " << i << ": "
        << (dist[i] == LLONG_MAX ? "INFINITY" : to_string(dist[i])) << "\n";

return 0;
}

```

6.4 Output




Figure 1: Graph

```

5 7
0 1 4
0 2 2
1 2 5
1 3 10
2 3 3
2 4 2
3 4 4
0
Shortest distances from node 0:
Node 0: 0
Node 1: 4
Node 2: 2
Node 3: 5
Node 4: 4

```